

Cloud Computing and Big Data

Hadoop and Map-Reduce

Oxford University
Software Engineering
Programme
July 2020



Contents

- Understanding Map Reduce
 - Functional programming patterns applied for scalability
- Hadoop
 - Map-reduce in Hadoop
 - Python
 - Java
 - HDFS
 - Yarn
 - Pig and Hive
- Further reading



Original 2008 Google Paper

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value pairs, and a *reduce* function that merges all intermediate values associated with the same intermediate key. Many real world tasks are expressible in this model, as shown

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to parallelize the computation, distribute the data, and handle failures conspire to obscure the original simple computation with large amounts of complex code to deal with these issues.



Yahoo 2007

Map-Reduce-Merge: Simplified Relational Data Processing on Large Clusters

Hung-chih Yang, Ali Dasdan
Yahoo!
Sunnyvale, CA, USA
{hcyang,dasdan}@yahoo-inc.com

Ruey-Lung Hsiao, D. Stott Parker
Computer Science Department, UCLA
Los Angeles, CA, USA
{rlhsiao,stott}@cs.ucla.edu



Programming Model

- Input & Output: each a set of key/value pairs
 - map (in_key, in_value) -> list(out_key, intermediate_value)
 - Processes input key/value pair
 - Produces set of intermediate pairs
 - reduce (out_key, list(intermediate_value)) -> list(out_value)
 - Combines all intermediate values for a particular key
 - Produces a set of merged output values (usually just one)
- Inspired by similar primitives in LISP and other languages

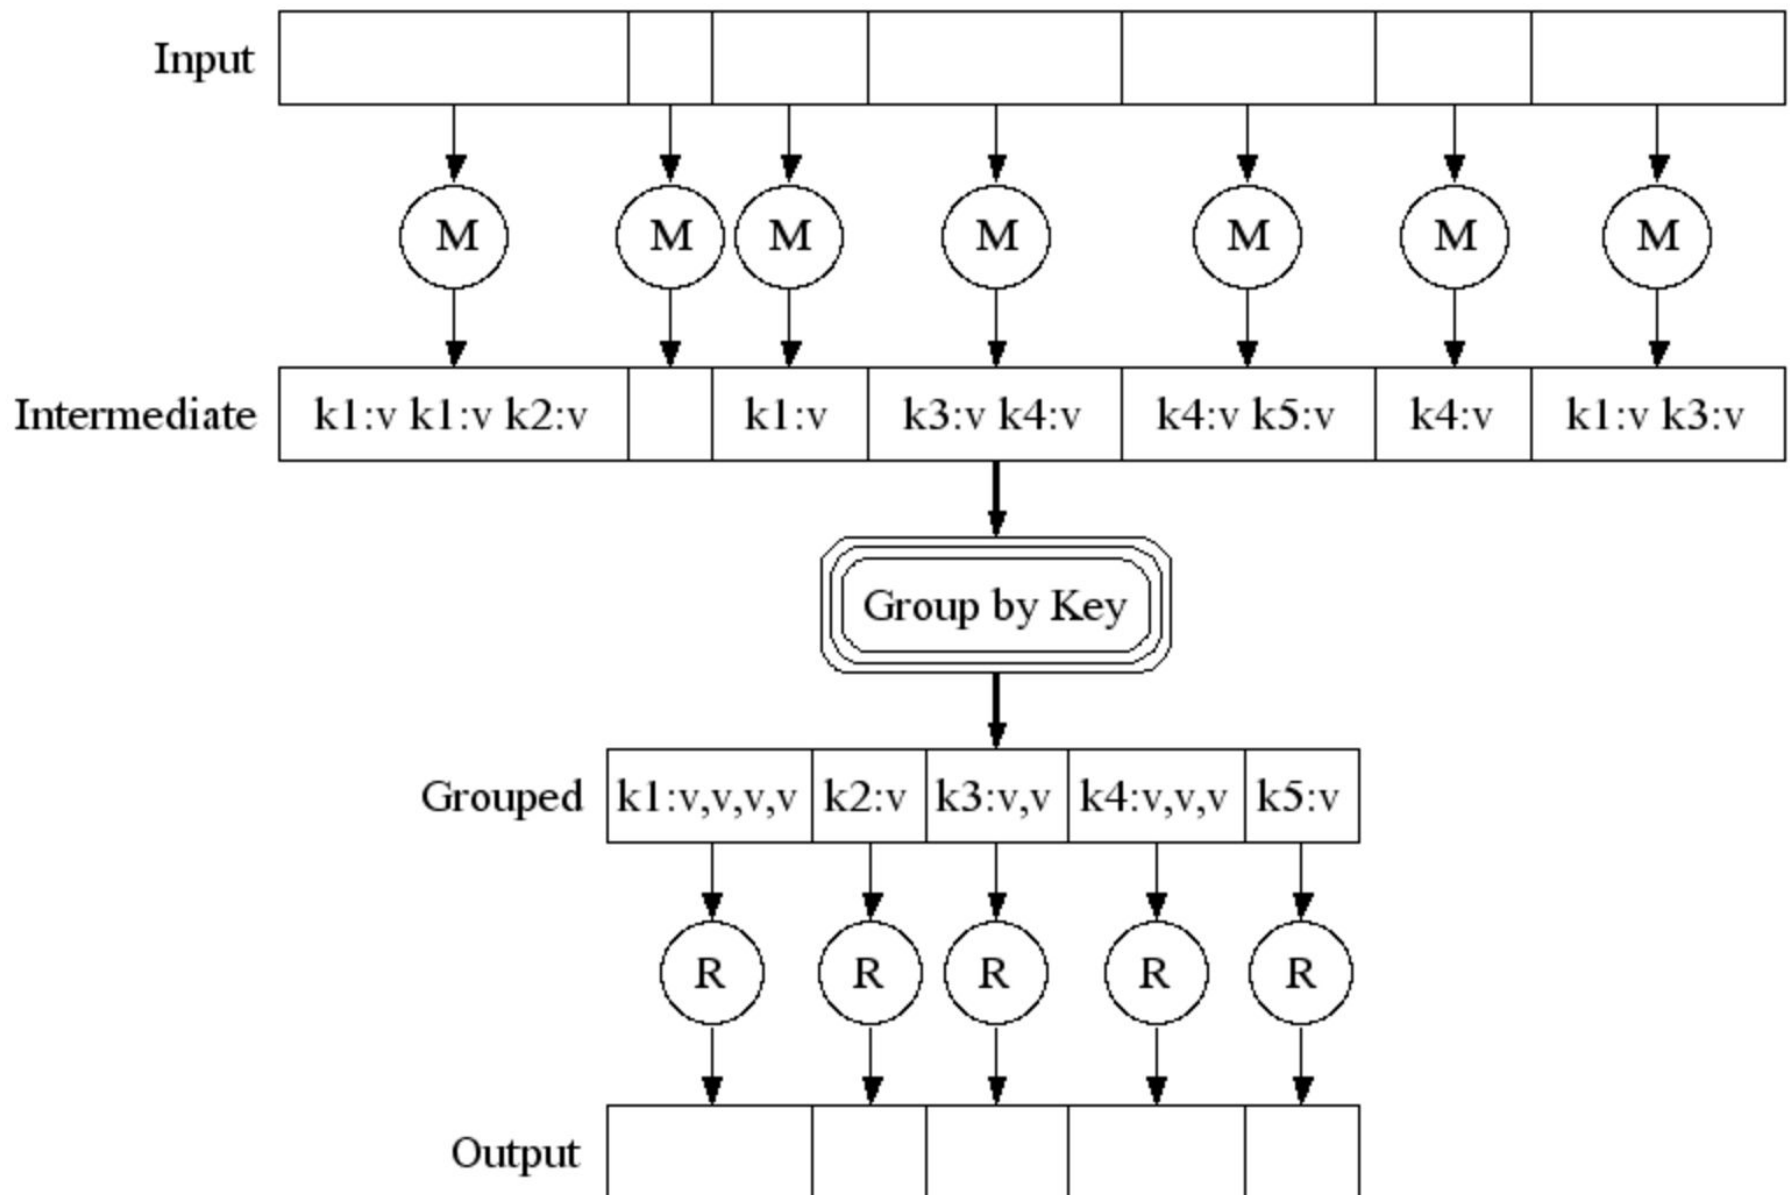


Class Exercise

- Find a post it note and write your name and day/month of birth on it.
- You don't need to divulge the year!


Paul
5
December





Google's early use of MR

Map Reduce programs in their code repository



Map Reduce example #1

Java

```
// some input
Integer array[] = {1,2,3,4,5};
// make a Stream of it
Stream<Integer> stream =
    Arrays.stream(array);
// map then reduce
Integer map_reduced =
    stream
        .map(x -> x*x)
        .reduce(0, (a, b)->a+b);
// should print 55
System.out.println(map_reduced);
```



Notice how

- We could have squared each number in the stream ***at the same time***
- Adding them up needed all the results available



Map Reduce example #2

in words

- Do a word count on 1000 books:
 - First count each book (**map** wc onto book)
 - Then **reduce** the outputs to a global wordcount across all books

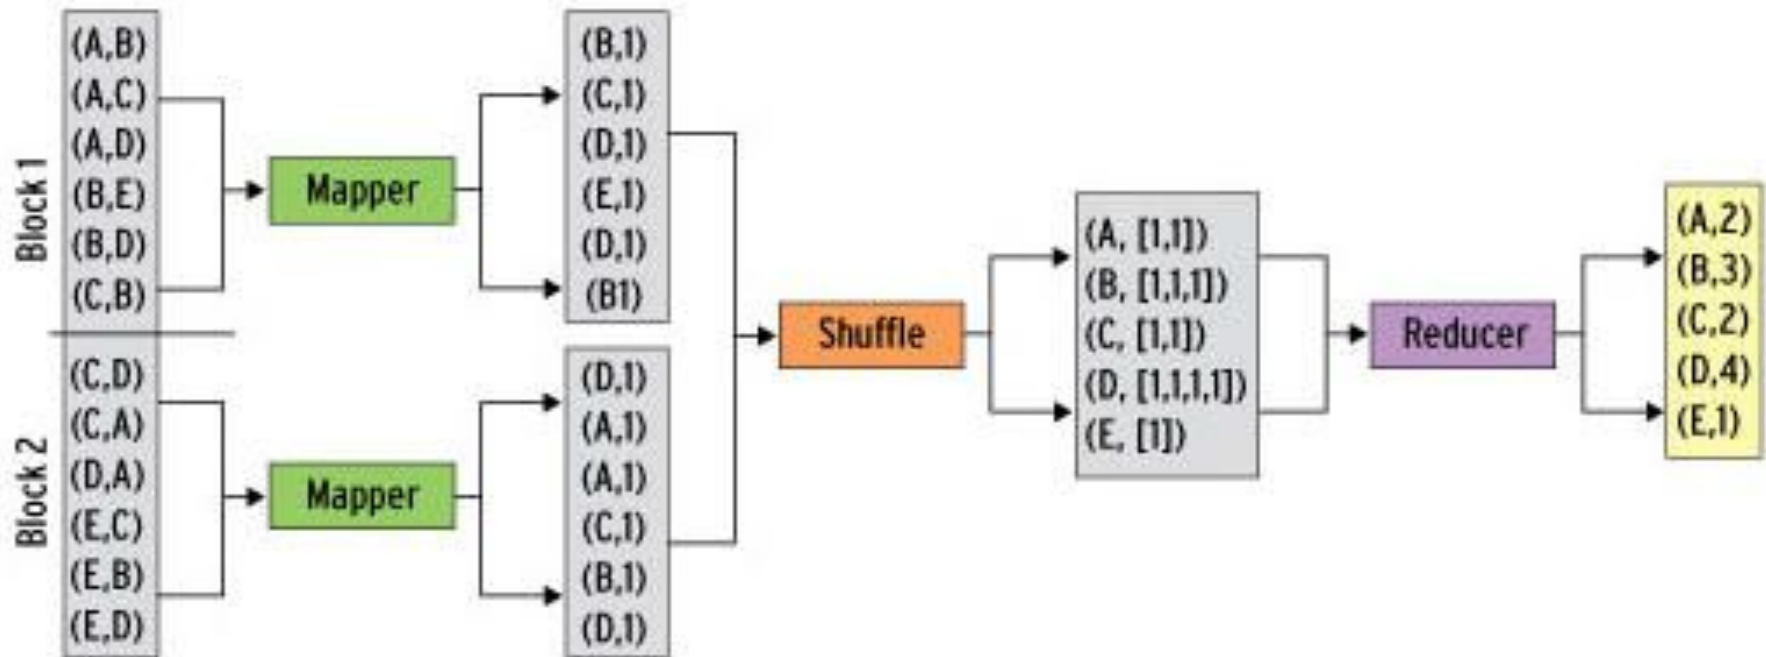


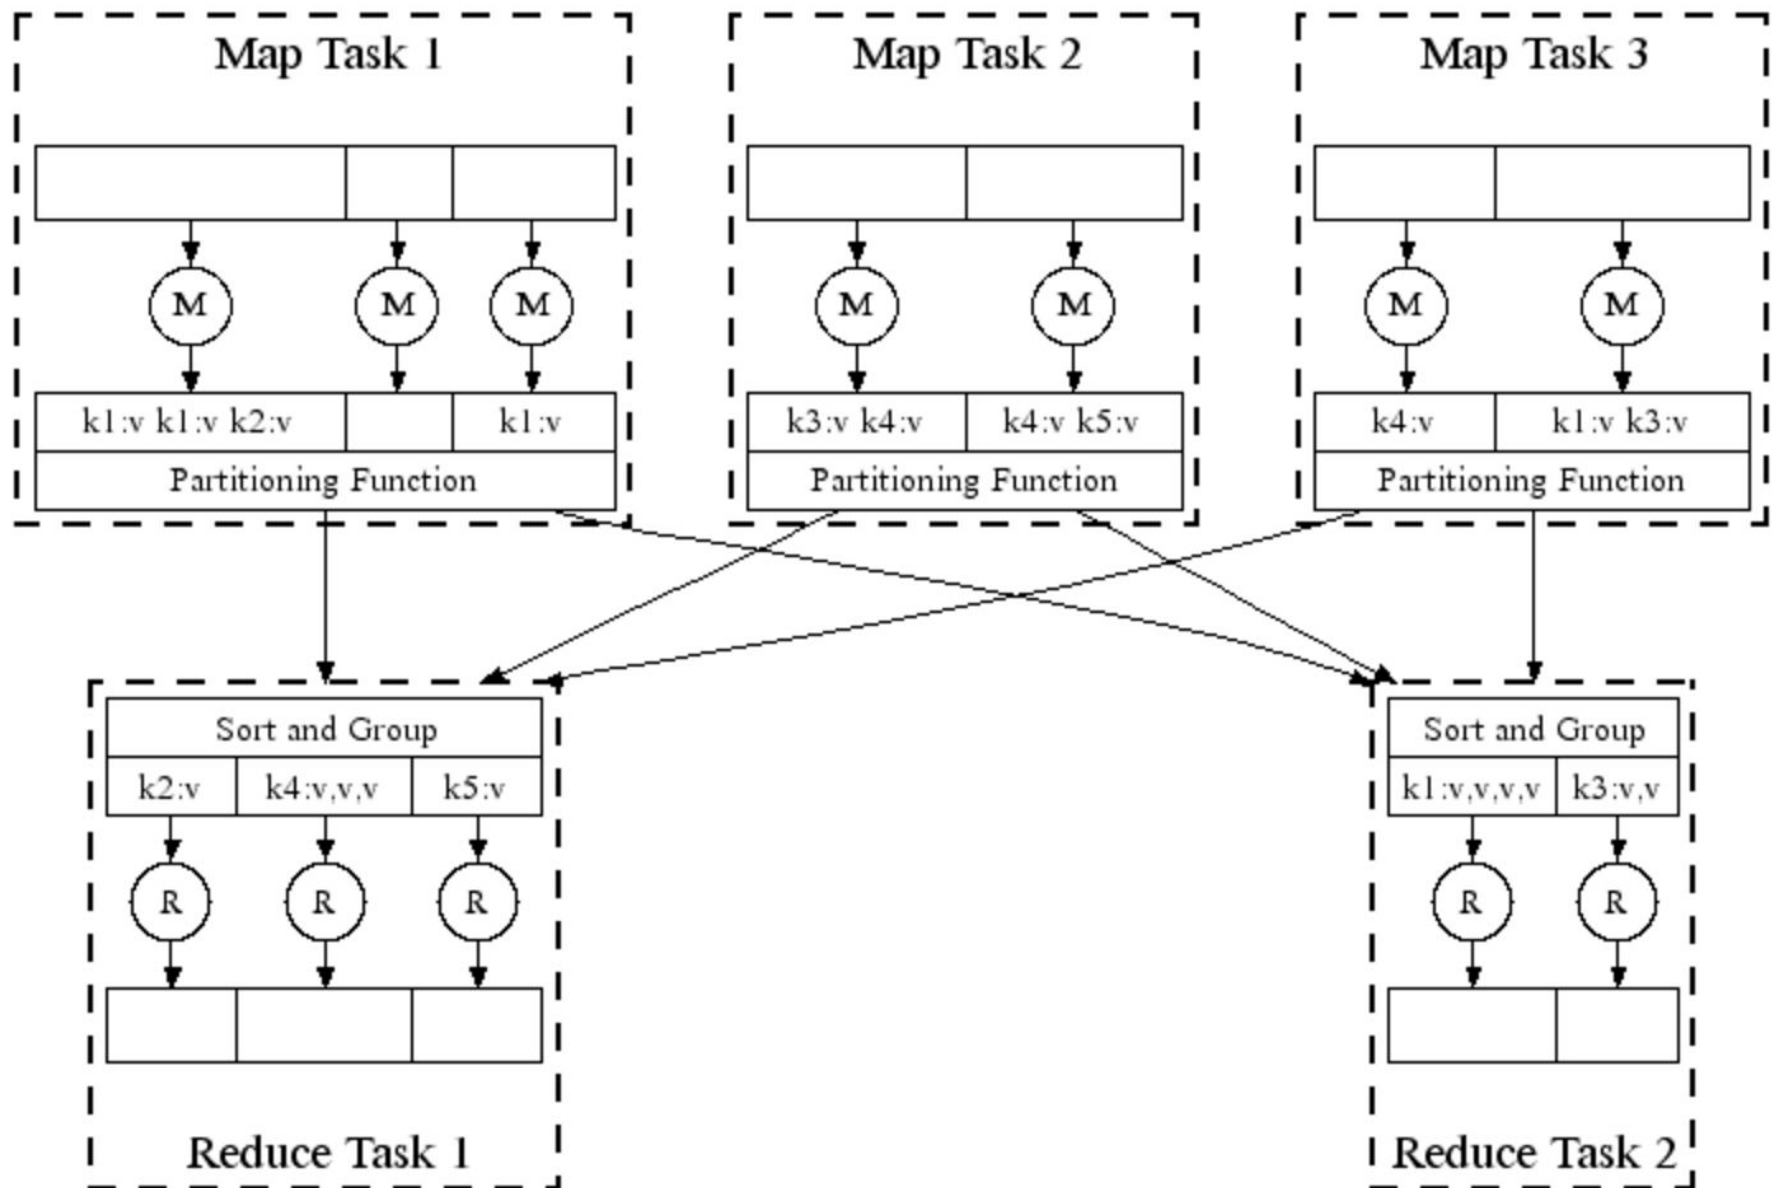
Efficiency

- Reduce phase:
 - We can theoretically process each word in parallel
- How?
 - Shuffle / Sort the results from the map phase by key (word)
 - Partition by keys
 - Parallelize the reduce phase



Map/Shuffle/Reduce





Map Reduce in Real Life

- Analysing web logs
 - Summarise by user / cookie
 - Then aggregate to identify who did what
- Analysing twitter data
 - Who retweeted
 - Who was retweeted the most
- Almost all big data problems can be re-factored into Map Reduce
 - Some more efficiently than others



Tuning

- Fault tolerance
 - Simply re-execute work that fails
- Performance:
 - Partitioning the data
 - Moving the work to near the data



Apache Hadoop

- The most famous and popular Map Reduce framework
 - Open Source
 - Written in Java, but supports other languages
 - Runs Map Reduce workloads across a cloud or cluster of machines
 - Supports a distributed filesystem to store data for these jobs
 - Provides reliability when servers in the cluster fail



Components of Hadoop

Map Reduce or Other Workloads

Java, Scala, Python, Apache Pig, Apache Hive, etc


YARN (Yet Another Resource Negotiator)
Cluster Resource Management

Hadoop Distributed File System (HDFS)

Redundant Reliable Distributed File System



HDFS in a nutshell



HDFS inspiration

- Google File System

- Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. 2003. The Google file system. In Proceedings of the nineteenth ACM symposium on Operating systems principles (SOSP '03). ACM, New York, NY, USA, 29-43.

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung
Google*

ABSTRACT

We have designed and implemented the Google File System, a scalable distributed file system for large distributed data-intensive applications. It provides fault tolerance while running on inexpensive commodity hardware, and it delivers high aggregate performance to a large number of clients.

While sharing many of the same goals as previous distributed file systems, our design has been driven by observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier file system assumptions. This has led us to reexamine traditional choices and explore radically different design points.

The file system has successfully met our storage needs. It is widely deployed within Google as the storage platform

1. INTRODUCTION

We have designed and implemented the Google File System (GFS) to meet the rapidly growing demands of Google's data processing needs. GFS shares many of the same goals as previous distributed file systems such as performance, scalability, reliability, and availability. However, its design has been driven by key observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier file system design assumptions. We have reexamined traditional choices and explored radically different points in the design space.

First, component failures are the norm rather than the exception. The file system consists of hundreds or even thousands of storage machines built from inexpensive com-



HDFS overview

- Good for streaming access to large files, reliability, scale
- Not good for random access, small files
- Blocks of data 64Mb in size (configurable)
- Each block can be replicated across multiple data nodes for High Availability (HA)



HDFS commands

- start-dfs.sh
- stop-dfs.sh
- `hadoop fs <command>`
 - e.g. `hadoop fs cat /user/hduser/file`
 - `hadoop fs mkdir -p /user/hduser/`
 - `hadoop fs put localfile /user/hduser/remotefile`
 - `hadoop fs get /user/hduser/remotefile localfile`



HDFS Usage

- Spotify has 1600+ nodes, storing 60+ petabytes of data
 - <https://www.usenix.org/system/files/conference/fast17/fast17-niazi.pdf>
- One of the Facebook's largest clusters (based on HDFS) holds more than 100 PB of data, processing more than 60,000 Hive queries a day
 - <https://www.facebook.com/notes/facebook-engineering/under-the-hood-scheduling-mapreduce-jobs-more-efficiently-with-corona/>



HopFS

- HopFS is a drop-in replacement for HDFS, based on HDFS v2.0.4.

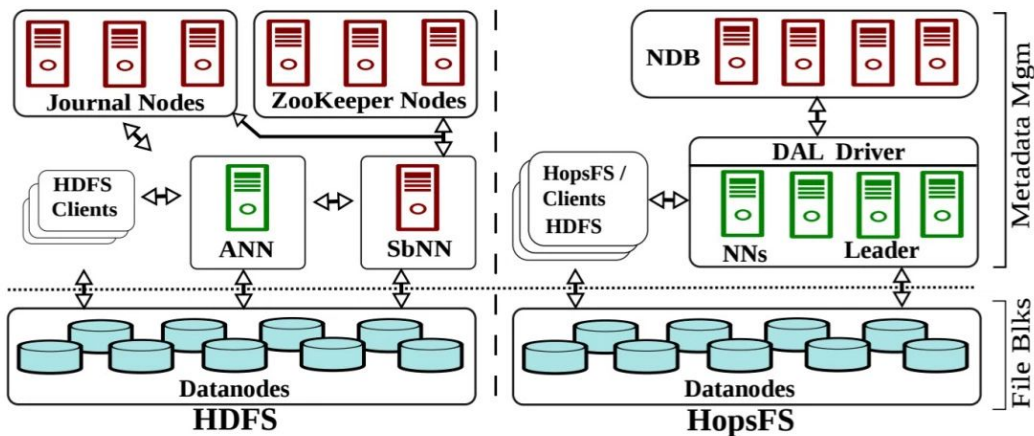


Figure 1: System architecture for HDFS and HopFS. For high availability, HDFS requires an Active NameNode (ANN), at least one Standby NameNode (SbNN), at least three Journal Nodes for quorum-based replication of the write ahead log of metadata changes, and at least three ZooKeeper instances for quorum based coordination. HopFS supports multiple stateless namenodes that access the metadata stored in NDB database nodes.





Figure 6: HopsFS and HDFS throughput for Spotify workload.

What is YARN?

- **YARN is the system that runs your code on multiple nodes**
- Hadoop 2.0 replacement for the cluster manager
 - Basically a model to distribute and manage workloads
 - Not just MapReduce but supports other workloads



YARN architecture



Map Reduce in Hadoop

- Map reduce in Hadoop actually consists of multiple steps
 - Mapper
 - Works on a single file, line by line
 - Combiner
 - Like a reducer, but still on a single system taking the output of the mapper
 - Reducer
 - Takes the outputs of multiple mapper/combiners



The general flow

- On the whole, we expect to produce key-value $\langle K, V \rangle$ pairs from any mapper or reducer
 - In some cases we may produce $\langle K, \langle V1, V2, \dots \rangle \rangle$
- The results are stored to file and then read from file



Mappers

- Take input files and produce <K,V> pairs
- Each mapper gets a complete file
- Each mapper runs on a single system



Combiners

- The combiner function must take the $\langle K, V \rangle$ output of the mapper
- Produce the same format $\langle K, V \rangle$
- Must be associative and commutative
- *Runs on the same node as the mapper*
- May actually be the reducer function if the reducer follows the rules above



Reducers

- The reducers get the $\langle K, V \rangle$ pairs output from the mappers/combiners
- The output is first sorted/partitioned
- The reducers produce the final expected output, usually in the form of $\langle K, V \rangle$ pairs



Summary

- Understanding the Map Reduce Model
- How is it implemented in Hadoop
- HDFS
- Yarn



Questions?



© Paul Fremantle 2015. This work is licensed under a Creative Commons
Attribution-NonCommercial-ShareAlike 4.0 International License
See <http://creativecommons.org/licenses/by-nc-sa/4.0/>